Hadoop Multi-Nodes Cluster Setup on Ubuntu 12.04

This configuration is referred to:

http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-multi-node-cluster/

1. Prerequisites

It is assumed that all the nodes have successfully configured as single-node cluster. To see how to configure hadoop single node cluster, refer the following post:

http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/

2. Networking

Edit machine name to master, slave1, slave2, slave3, slave4, slave6 and make sure any of them can reach others over the network

	On master node: change host name to master on master node at: /etc/hostname
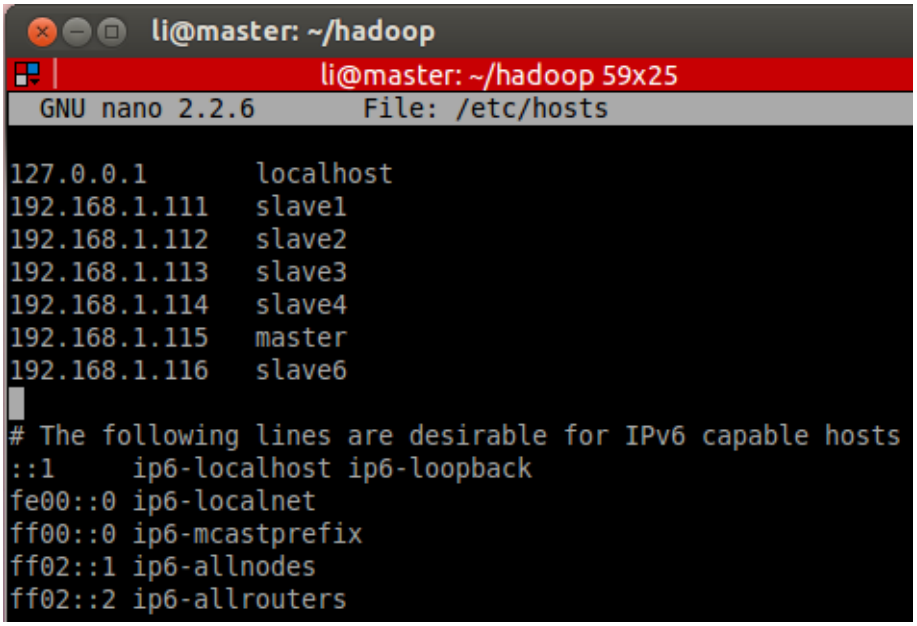	On master node: Add master and slave machine IP & name
	On slave1: change host name to slave1 on slave1 node at:

	/etc/hostname
li@slave1: ~/hadoop/logs 64x18 GNU nano 2.2.6 File: /etc/hosts 127.0.0.1 localhost 192.168.1.111 slave1 192.168.1.112 slave2 192.168.1.113 slave3 192.168.1.114 slave4 192.168.1.115 master 192.168.1.116 slave6 # The following lines are desirable for IPv6 capable hosts ::1 ip6-localhost ip6-loopback fe00::0 ip6-localnet ff00::0 ip6-mcastprefix ff02::1 ip6-allnodes	On slave1 node: Add master and slave machine IP & name
Repeat the same steps for slave2, 3, 4, 6	

3. SSH Access

Enable user li on master to connect to slave1, 2, 3, 4, 6 via a password-less SSH login

li@master: ~ 90x25 li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave1	On master node: add the li@master's public SSH key (which should be in $HOME/.ssh/id_rsa.pub) to the authorized_keys file of li@slave1
li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave1 li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave2 li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave3 li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave4 li@master:~$ ssh-copy-id -i $HOME/.ssh/id_rsa.pub li@slave6	Repeat the same step to add SSH key to slave2, 3, 4, 6

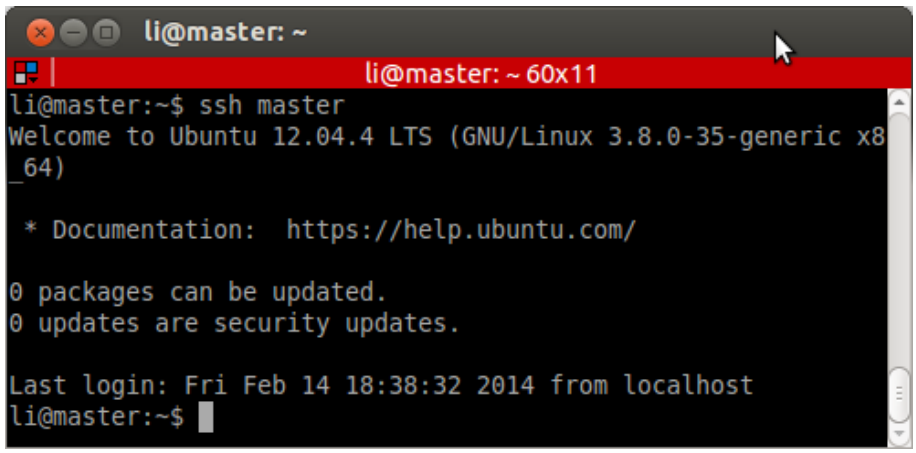

Test

on master:
connect from master to master

ssh master

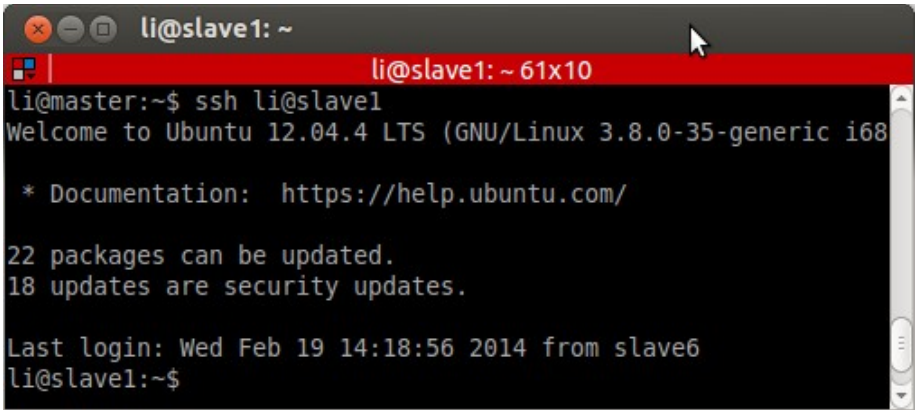

Test

on master

connect from master to slave1

ssh li@slave1

5. Configuration - Hadoop/conf/masters(master only)

Despite its name, the conf/masters file defines on which machines Hadoop will start *secondary NameNodes* in our multi-node cluster. In our case, this is just the master machine. The primary NameNode and the JobTracker will be started on the same machine if you run bin/start-all.sh.

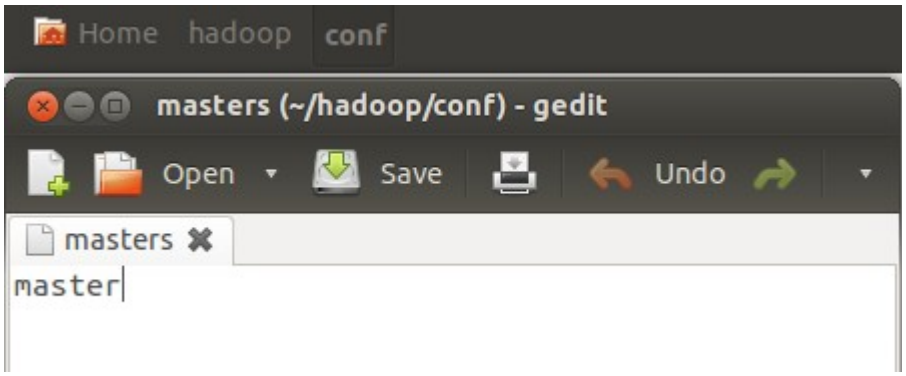

Edit hadoop/conf/masters:

Hadoop/conf/slaves(master only)

The conf/slaves file lists the hosts, one per line, where the Hadoop slave daemons (DataNodes and TaskTrackers) will be run. We want both the master box and the slave box to act as Hadoop slaves because we want both of them to store and process data.

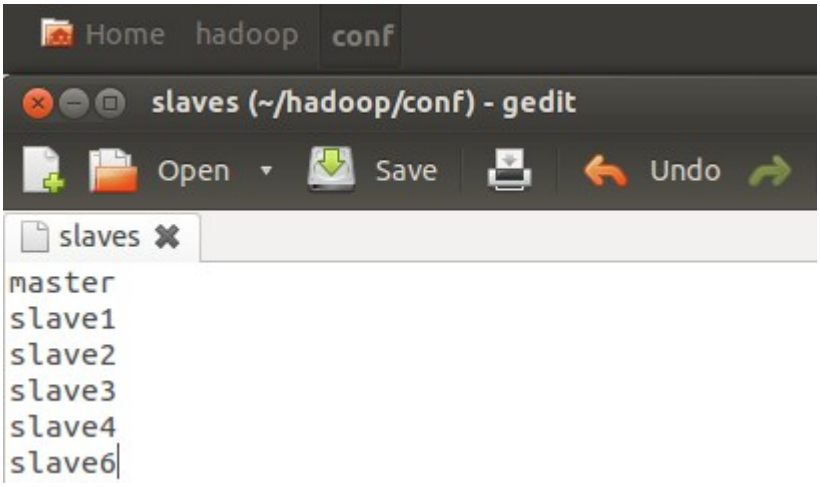	Edit hadoop/conf/slaves:

conf/*-site.xml (all machines)

First, we have to change the fs.default.name parameter (in conf/core-site.xml), which specifies the NameNode (the HDFS master) host and port. In our case, this is the master machine.

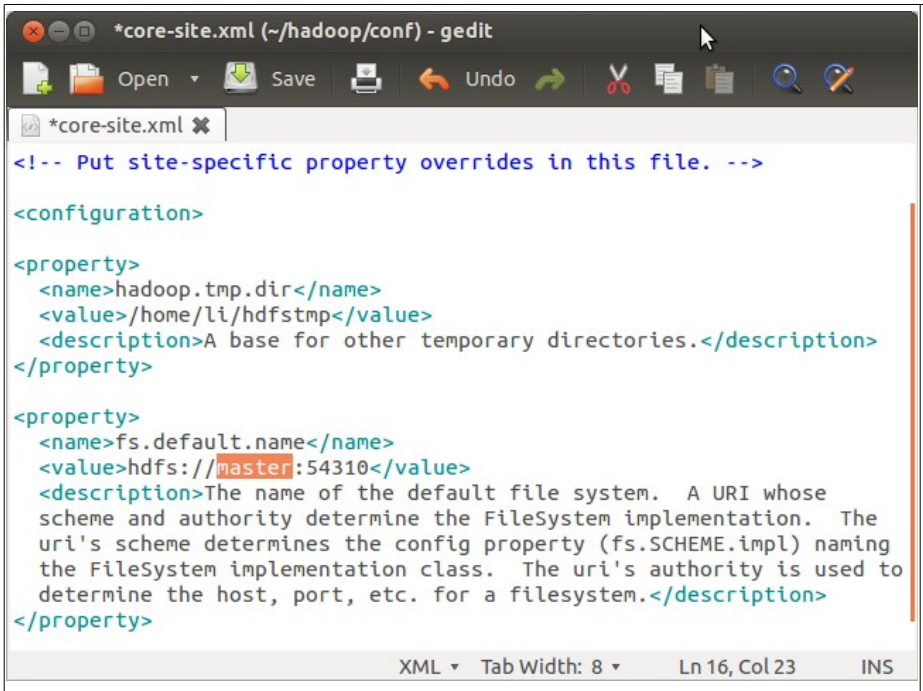	Edit hadoop/conf/core-site.xml

Second, we have to change the mapred.job.tracker parameter (in conf/mapred-site.xml), which specifies the JobTracker (MapReduce master) host and port. Again, this is the master inour case.

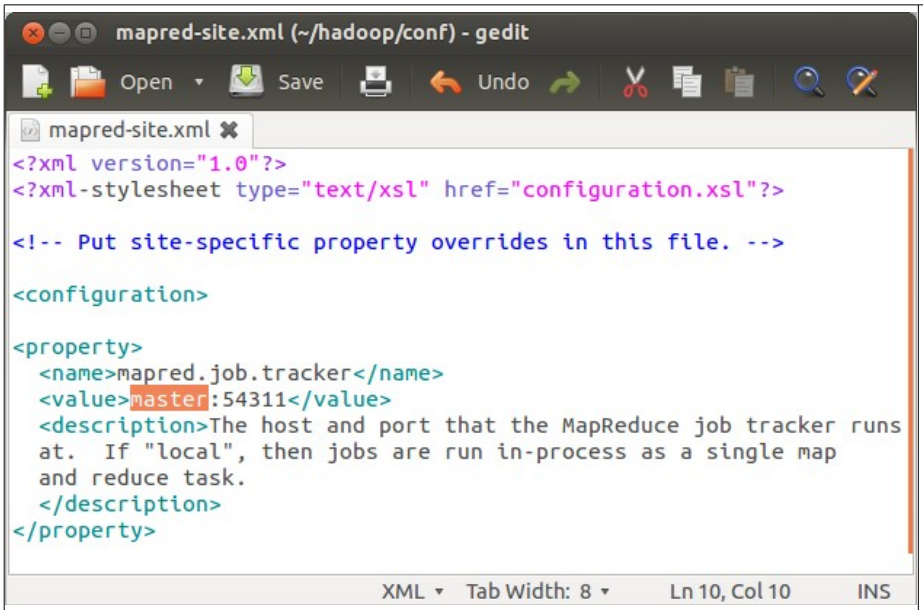

Edit hadoop/conf/mapred-site.xml

Third, we change the dfs.replication parameter (in conf/hdfs-site.xml) which specifies the default block replication. It defines how many machines a single file should be replicated to before it becomes available. If you set this to a value higher than the number of available slave nodes (more precisely, the number of DataNodes), you will start seeing a lot of “(Zero targets found, forbidden1.size=1)” type errors in the log files.

The default value of dfs.replication is 3.

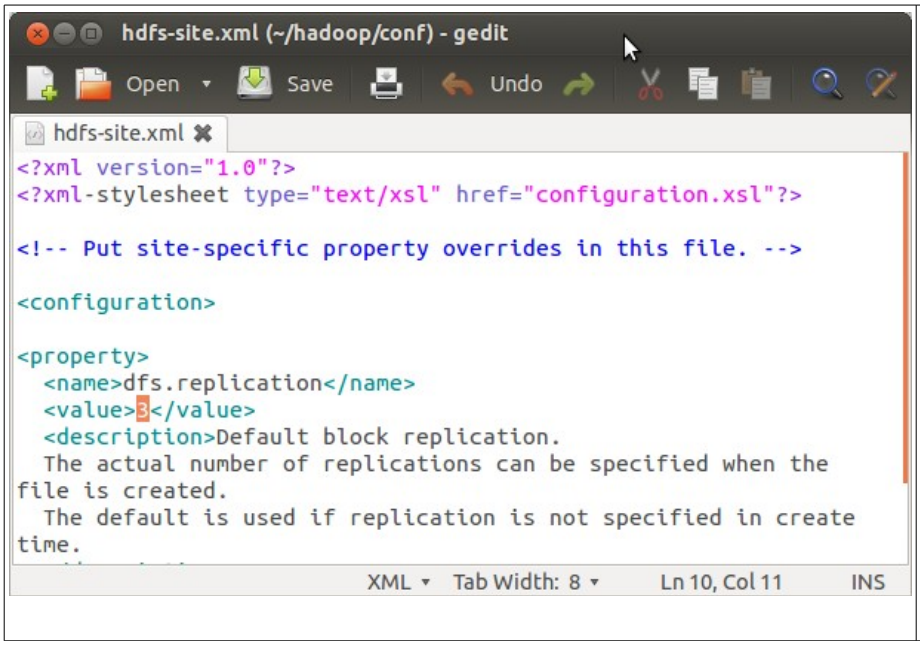

Edit hadoop/conf/hdfs-site.xml

6. Formatting the HDFS filesystem via the NameNode

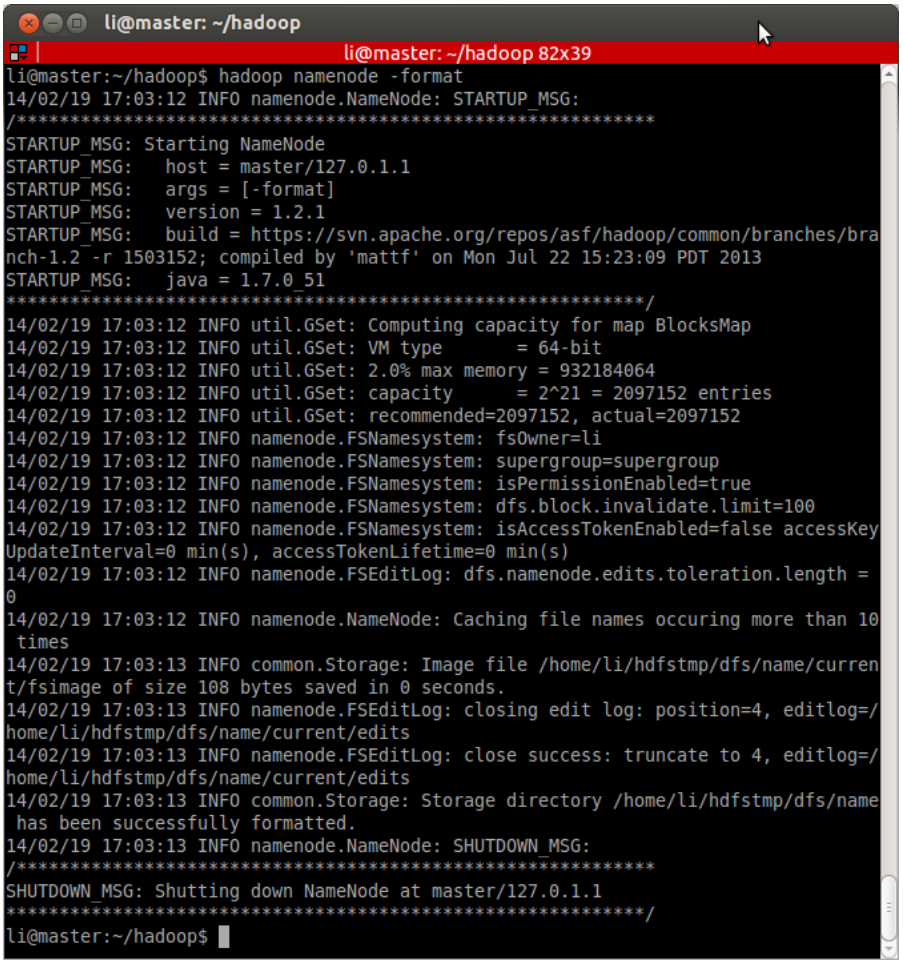

hadoop namenode -format

7. Starting Check and stop the multi-node cluster

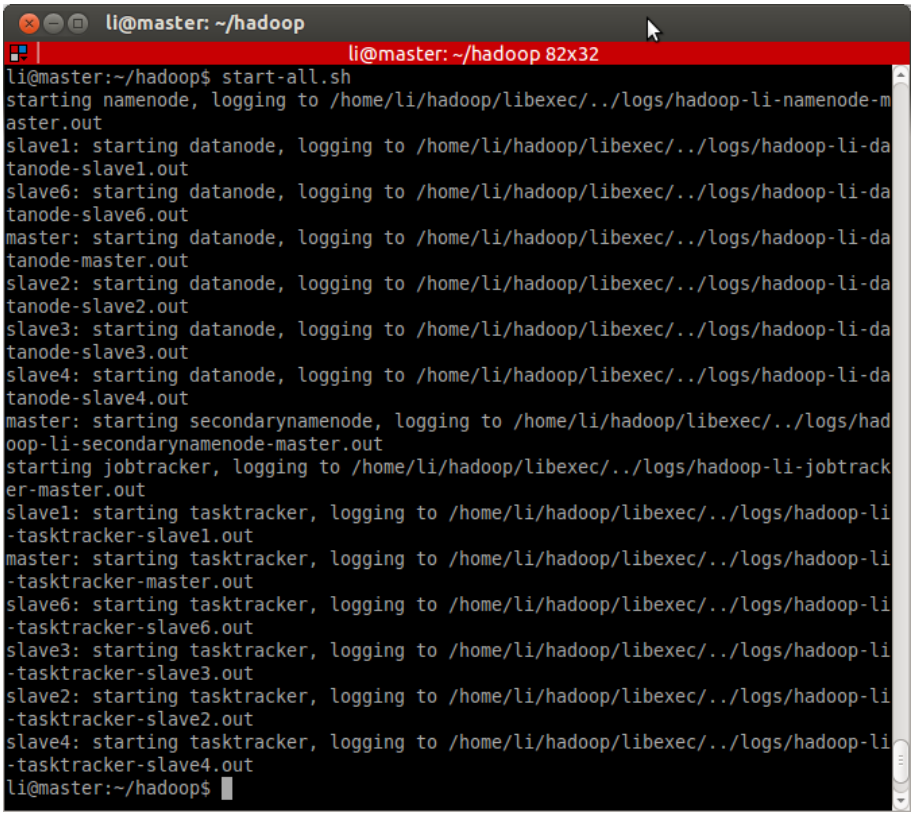

To start the cluster:

on master:

start-all.sh

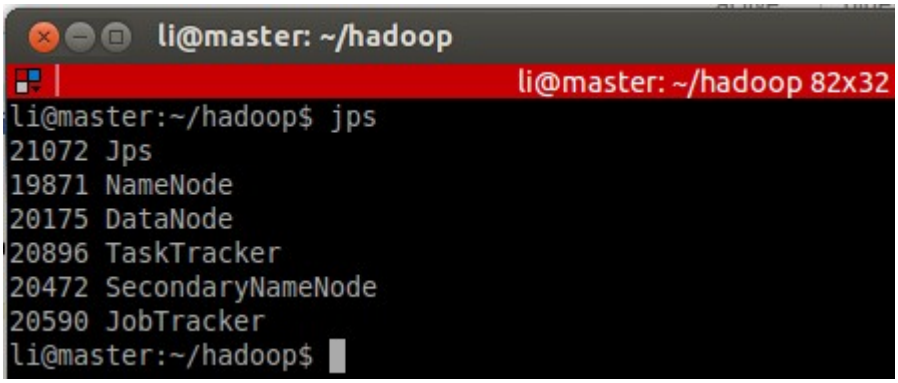

Check running java processes on master:

jps

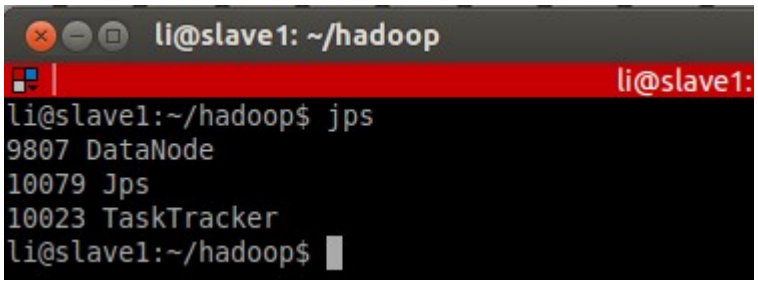

Check running java processes on slave1:

jps

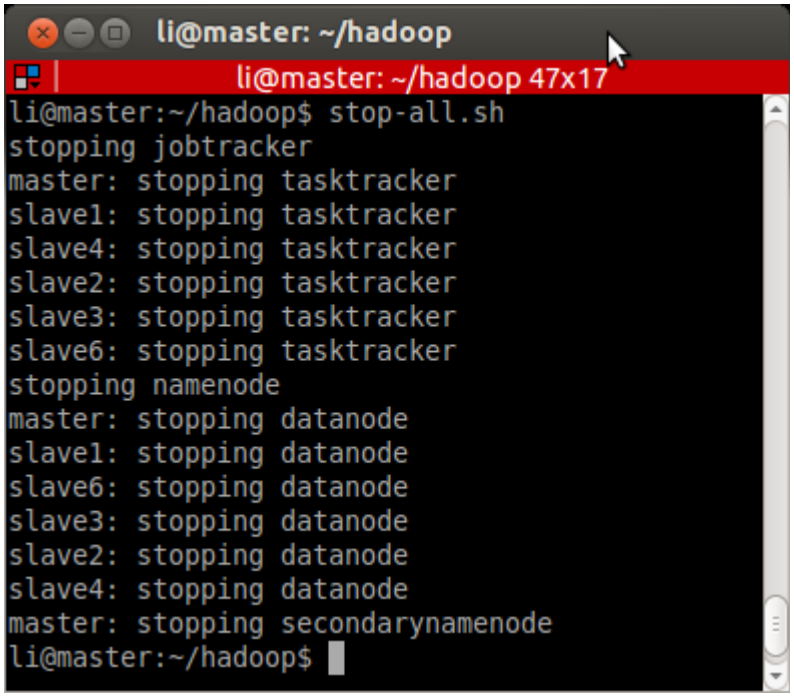

Stop cluster

on master:

stop-all.sh

8. Troubleshooting

1) If datanode does not run on master after start the cluster, delete the hdfstmp folder and create a new empty hdfstmp folder then run

hadoop namenode -format

2) If you got connection refused error like

ERROR security.UserGroupInformation: PriviledgedActionException as:li cause:java.net.ConnectException: Call to master/192.168.1.115:54311 failed on connection exception: java.net.ConnectException: Connection refused

Try let the namenode leave safemode:

```
bin/hadoop dfsadmin -safemode leave
```